# ADITYA ANAND

adianand2002@gmail.com | +91 8728009774 | linkedin.com/in/adianand912 | github.com/Adiii1436

## WORK EXPERIENCE

**INFOGAIN CORPORATION** Bangalore, India
Associate Software Engineer July 2025 – Present

- Designed and operated distributed PySpark ETL pipelines on Azure Databricks, processing GBs of enterprise data across a multi-layer cloud architecture.
- Implemented automated CDC reconciliation frameworks to continuously validate data integrity across distributed systems, embedding a test-driven, CI/CD-aligned engineering discipline into the pipeline lifecycle.
- Collaborated with project managers and technical leads to translate business requirements into compliant, production-ready ingestion pipelines, practising Agile delivery and iterative stakeholder feedback cycles.

**SAMSUNG RESEARCH AND DEVELOPMENT** Bangalore, India (Remote)
Research Intern Feb 2023 – Sep 2023

- Designed a multi-node distributed microservices system on Kubernetes, cutting API response time by 25% and server load by 30%.
- Reduced log retrieval time by 30% by integrating an Open Log Forwarder into existing Kubernetes infrastructure, demonstrating ability to extend and improve large-scale, production systems with agility.
- Accelerated peak-hour troubleshooting by 25% by building real-time Grafana dashboards for system observability - critical for monitoring inference services at scale.

## PROJECTS

**Alloy AI - Code Assistant |** *TypeScript, Node.js, VS Code API*

- Architected a privacy-first AI coding assistant using RAG: a local Merkle Tree Indexer maps code dependencies to enable high-precision context retrieval without ever uploading the user's codebase.
- Shipped to the VS Code Marketplace with multi-LLM backend support (Gemini, OpenAI), and built a custom token compression algorithm to handle large error logs efficiently.

**Context Aware File Organizer |** *Python, Llama.cpp, Pyside6*

- Built a fully offline generative AI pipeline that semantically clusters documents using local vector embeddings and Agglomerative Clustering, then feeds cluster summaries into a locally-running Llama inference engine to generate context-aware folder names.
- Chose llama.cpp as the inference runtime deliberately for its quantization support and CPU/GPU portability, gaining hands-on understanding of the LLM inference stack that underpins services like AOAI.

**MLOps Pipeline – CNN on FashionMNIST |** Python, PyTorch, ZenML, MLflow, DagsHub

- Built an end-to-end MLOps pipeline covering data ingestion, CNN training, evaluation, and automated deployment.
- Integrated MLflow experiment tracking with a deployment trigger gating releases on cross-entropy loss thresholds, applying the same CI/CD discipline used in real-world model serving systems.

## SKILLS

**Languages**: Python, C, C++, SQL.
**AI / GenAI:** Pytorch, Langchain, Chromadb, RAG, Transformer architectures, LLM inference optimization.
**Distributed Systems:** Kubernetes, Docker, Apache Spark (PySpark), Microservices.
**Cloud:** Azure Databricks, Google Cloud Platform (GCP).
**DevOps:** Git, Linux, Grafana, Elasticsearch, CI/CD.

## CERTIFICATIONS

- Databricks Data Engineer Associate.
- Google Cloud Certified Professional Data Engineer.

## EDUCATION

**CHANDIGARH UNIVERSITY** Chandigarh, India
Bachelor of Engineering Jun 2021 - Jun 2025
Major in Computer Science with Specialization in AIML
CGPA: 7.8/10
Relevant Coursework: Artificial Intelligence; Machine Learning; Distributed Systems; Data Structure & Algorithms; Computer Networks.